

CS 455 – Spring 2018

RUNNING THE WORD COUNT EXAMPLE

Before starting, make sure that you have HDFS and Yarn running, using `sbin/start-dfs.sh` and `sbin/start-yarn.sh`

- Download text copies of at least 3 books from Project Gutenberg: (<http://www.gutenberg.org/>)

```
st-vrain> la
total 1284
-rw----- 1 class 84358 Mar 26 23:47 Much_Ado_About_Nothing.txt
-rw----- 1 class 792920 Mar 26 23:48 Tale_Of_Two_Cities.txt
-rw----- 1 class 421884 Mar 26 23:49 Tom_Sawyer.txt
```

- Create a directory in your local space on HDFS to store these books:

```
st-vrain> $HADOOP_HOME/bin/hdfs dfs -mkdir /cs455
st-vrain> $HADOOP_HOME/bin/hdfs dfs -mkdir /cs455/books
st-vrain> $HADOOP_HOME/bin/hdfs dfs -ls /cs455/
Found 1 items
drwxr-xr-x - cs455 supergroup          0 2018-03-08 23:51 /cs455/books
```

- Move the books from NFS into HDFS:

```
st-vrain> $HADOOP_HOME/bin/hdfs dfs -put *.txt /cs455/books
st-vrain> $HADOOP_HOME/bin/hdfs dfs -ls /cs455/books
Found 3 items
-rw-r--r-- 3 cs455 supergroup      84358 2018-03-08 23:55
/cs455/books/Much_Ado_About_Nothing.txt
-rw-r--r-- 3 cs455 supergroup     792920 2018-03-08 23:55
/cs455/books/Tale_Of_Two_Cities.txt
-rw-r--r-- 3 cs455 supergroup     421884 2018-03-08 23:55
/cs455/books/Tom_Sawyer.txt
```

- You can also check that the books are there via the HDFS web portal:

Hadoop Overview Datanodes Snapshot Startup Progress Utilities							
Browse Directory							
<input type="text" value="/cs455/books"/>							<input type="button" value="Go"/>
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	cs455	supergroup	457.84 KB	3/7/2018, 1:56:47 PM	7	128 MB	Paper-making.txt
-rw-r--r--	cs455	supergroup	370.4 KB	3/7/2018, 1:57:05 PM	7	128 MB	Salt-Lake-City.txt
-rw-r--r--	cs455	supergroup	414.85 KB	3/7/2018, 1:57:30 PM	7	128 MB	The-Last-Three-Soldiers.txt
Hadoop, 2016.							

- Download the source code of the word count example from CS 455 course web site. (link: <http://www.cs.colostate.edu/~cs455/cs455-wordcount-sp18.tar>)

```
wget http://www.cs.colostate.edu/~cs455/cs455-wordcount-sp18.tar
```

- Extract the tarball.

```
tar -xvf cs455-wordcount-sp18.tar
```

- This includes an Ant build file called build.xml. This is used to compile source and package it into a jar. After compiling, it will create the jar file inside the ./dist directory. You can use this build.xml file as it is for HW3-PC. Type 'ant' to compile the source and create the jar file.

```
st-vrain> ant
Buildfile: /s/bach/a/class/cs455/sp18-hadoop/cs455-wordcount-sp18/build.xml
init:
compile:
dist:
BUILD SUCCESSFUL
Total time: 0 seconds
st-vrain> ls ./dist/
wordcount.jar
st-vrain>
```

- Run the jar in yarn:

```
st-vrain> $HADOOP_HOME/bin/hadoop jar dist/wordcount.jar
cs455.hadoop.wordcount.WordCountJob /cs455/books
/cs455/wordcount-out
2017-03-08 00:36:53,833 INFO [main] client.RMProxy
(RMProxy.java:createRMProxy(98)) - Connecting to ResourceManager
at st-vrain/129.82.47.128:46783
2017-03-08 00:36:54,325 WARN [main] mapreduce.JobSubmitter
(JobSubmitter.java:copyAndConfigureFiles(153)) - Hadoop command-
```

```
line option parsing not performed. Implement the Tool interface
and execute your application with ToolRunner to remedy this.
2017-03-08 00:36:54,606 INFO [main] input.FileInputFormat
(FileInputFormat.java:listStatus(281)) - Total input paths to
process : 3
2017-03-08 00:36:54,696 INFO [main] mapreduce.JobSubmitter
(JobSubmitter.java:submitJobInternal(494)) - number of splits:3
2017-03-08 00:36:54,909 INFO [main] mapreduce.JobSubmitter
(JobSubmitter.java:printTokens(583)) - Submitting tokens for job:
job_1427438142863_0001
2017-03-08 00:36:55,231 INFO [main] impl.YarnClientImpl
(YarnClientImpl.java:submitApplication(251)) - Submitted
application application_1427438142863_0001
2017-03-08 00:36:55,270 INFO [main] mapreduce.Job
(Job.java:submit(1300)) - The url to track the job: http://st-
vrain:8088/proxy/application_1427438142863_0001/
2017-03-08 00:36:55,271 INFO [main] mapreduce.Job
(Job.java:monitorAndPrintJob(1345)) - Running job:
job_1427438142863_0001
2017-03-08 00:37:01,455 INFO [main] mapreduce.Job
(Job.java:monitorAndPrintJob(1366)) - Job job_1427438142863_0001
running in uber mode : false
2017-03-08 00:37:01,456 INFO [main] mapreduce.Job
(Job.java:monitorAndPrintJob(1373)) - map 0% reduce 0%
2017-03-08 00:37:07,530 INFO [main] mapreduce.Job
(Job.java:monitorAndPrintJob(1373)) - map 100% reduce 0%
2017-03-08 00:37:16,599 INFO [main] mapreduce.Job
(Job.java:monitorAndPrintJob(1373)) - map 100% reduce 100%
2017-03-08 00:37:17,631 INFO [main] mapreduce.Job
(Job.java:monitorAndPrintJob(1384)) - Job job_1427438142863_0001
completed successfully
2017-03-08 00:37:17,773 INFO [main] mapreduce.Job
(Job.java:monitorAndPrintJob(1391)) - Counters: 49
    File System Counters
        FILE: Number of bytes read=549269
        FILE: Number of bytes written=1522267
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=1299517
        HDFS: Number of bytes written=314863
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=3
        Launched reduce tasks=1
        Data-local map tasks=3
        Total time spent by all maps in occupied slots
(ms)=10788
```

```
Total time spent by all reduces in occupied slots
(ms)=6841
Total time spent by all map tasks (ms)=10788
Total time spent by all reduce tasks (ms)=6841
Total vcore-seconds taken by all map tasks=10788
Total vcore-seconds taken by all reduce tasks=6841
Total megabyte-seconds taken by all map tasks=11046912
Total megabyte-seconds taken by all reduce
tasks=7005184
Map-Reduce Framework

Map input records=27088
Map output records=226606
Map output bytes=2171352
Map output materialized bytes=549281
Input split bytes=355
Combine input records=226606
Combine output records=38119
Reduce input groups=29082
Reduce shuffle bytes=549281
Reduce input records=38119
Reduce output records=29082
Spilled Records=76238
Shuffled Maps =3
Failed Shuffles=0
Merged Map outputs=3
GC time elapsed (ms)=128
CPU time spent (ms)=8450
Physical memory (bytes) snapshot=956612608
Virtual memory (bytes) snapshot=3741761536
Total committed heap usage (bytes)=805306368

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=1299162
File Output Format Counters
Bytes Written=314863
```

- Check output in HDFS:

```
st-vrain> $HADOOP_HOME/bin/hdfs dfs -ls /cs455/wordcount-out
Found 2 items
-rw-r--r--  3 cs455 supergroup          0 2018-03-07 16:43
```

```
/cs455/wordcount-out/_SUCCESS
-rw-r--r--  3 cs455 supergroup    314863 2018-03-07 16:43
/cs455/wordcount-out/part-r-00000
```

Check the output in the web portal. Click on part-r-00000 file and it will be downloaded.

The screenshot shows the Hadoop web portal interface. At the top is a green navigation bar with links: Hadoop, Overview, Datanodes, Snapshot, Startup Progress, and Utilities. Below this is the 'Browse Directory' section. A search bar contains the path '/cs455/wordcount-out' and a 'Go!' button. Below the search bar is a table with the following columns: Permission, Owner, Group, Size, Last Modified, Replication, Block Size, and Name. The table contains two entries: one for '_SUCCESS' (0 B) and one for 'part-r-00000' (286.9 KB). Both files are owned by 'cs455' and 'supergroup'. At the bottom of the page, it says 'Hadoop, 2016.'

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	cs455	supergroup	0 B	3/7/2018, 4:43:53 PM	7	128 MB	_SUCCESS
-rw-r--r--	cs455	supergroup	286.9 KB	3/7/2018, 4:43:52 PM	7	128 MB	part-r-00000

- **NOTE:** if you run this repeatedly, you will need to either modify your output folder name, or delete it between runs
- Check the following link for the complete set of HDFS commands.
[<http://hadoop.apache.org/docs/current/hadoop-project-dist/hadoop-hdfs/HDFSCommands.html>]